

Hierarchical Opinion Classification using Large Language Models

FIRE 2025 — Forum for Information Retrieval Evaluation

Shuvam Banerji Seal

sbs22ms076@iiserkol.ac.in

Department of Chemical Sciences,
Indian Institute of Science Education & Research
Kolkata

Alok Mishra

maa24ms215@iiserkol.ac.in

Indian Institute of Science Education & Research
Kolkata

Utkarsha Ghosh

utkarsha.ghosh2023@iem.edu.in

Department of Information Technology,
Institute of Engineering and Management, Kolkata

December 17–20, 2025

1. Outline

Presentation Outline

- ▷ Outline
- ▷ Introduction
- ▷ Problem Statement
- ▷ Data Preparation
- ▷ Model Architecture
- ▷ Training Methodology
- ▷ Instruction Fine-Tuning
- ▷ Results
- ▷ Challenges & Solutions
- ▷ Conclusion
- ▷ Acknowledgments

2. Introduction

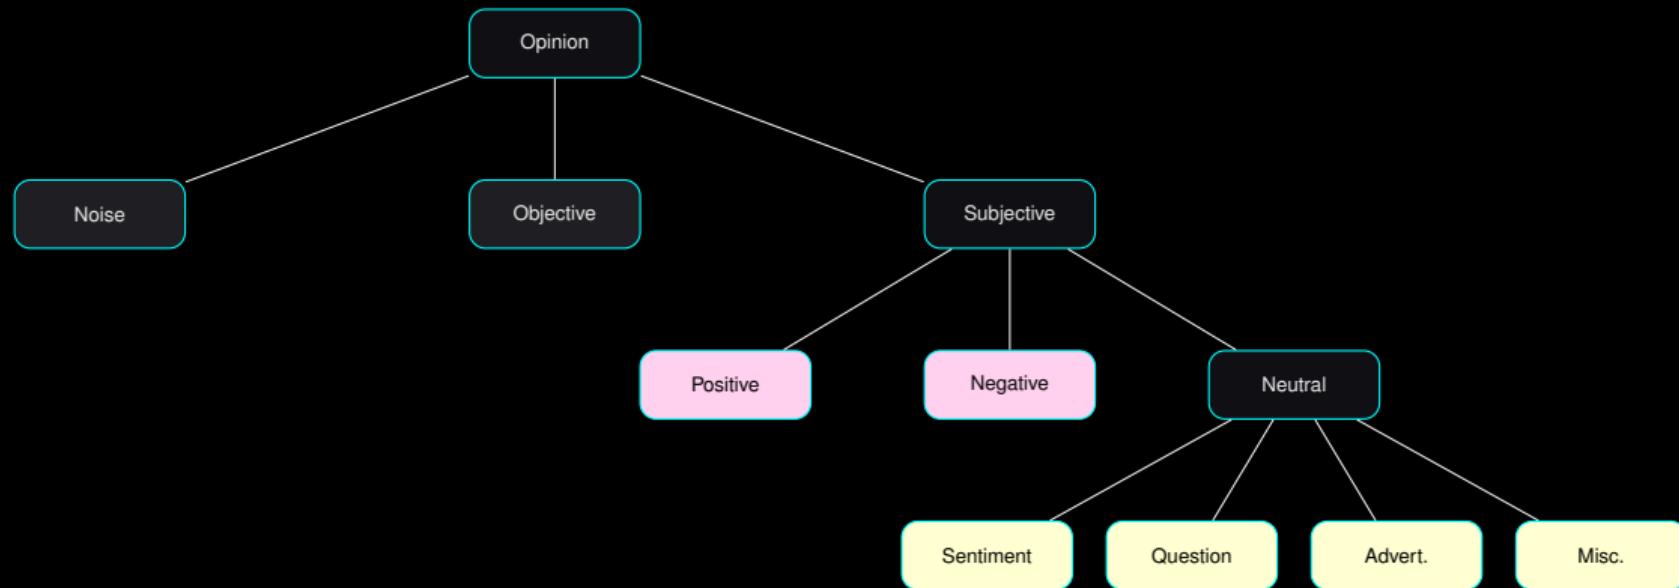
Context: FIRE 2025 Shared Task

- ◊ The datasets (Reddit, Twitter, YouTube) and the **3-level hierarchy** were provided as part of the shared task.
- ◊ **Why Hierarchy Matters:** It disambiguates *intent* from *sentiment*.
 - Differentiates **Complaints** (Negative) from **Inquiries** (Questions).
 - *Where it doesn't matter:* **Noise** is flat and requires no depth.

Real Examples from Dataset:

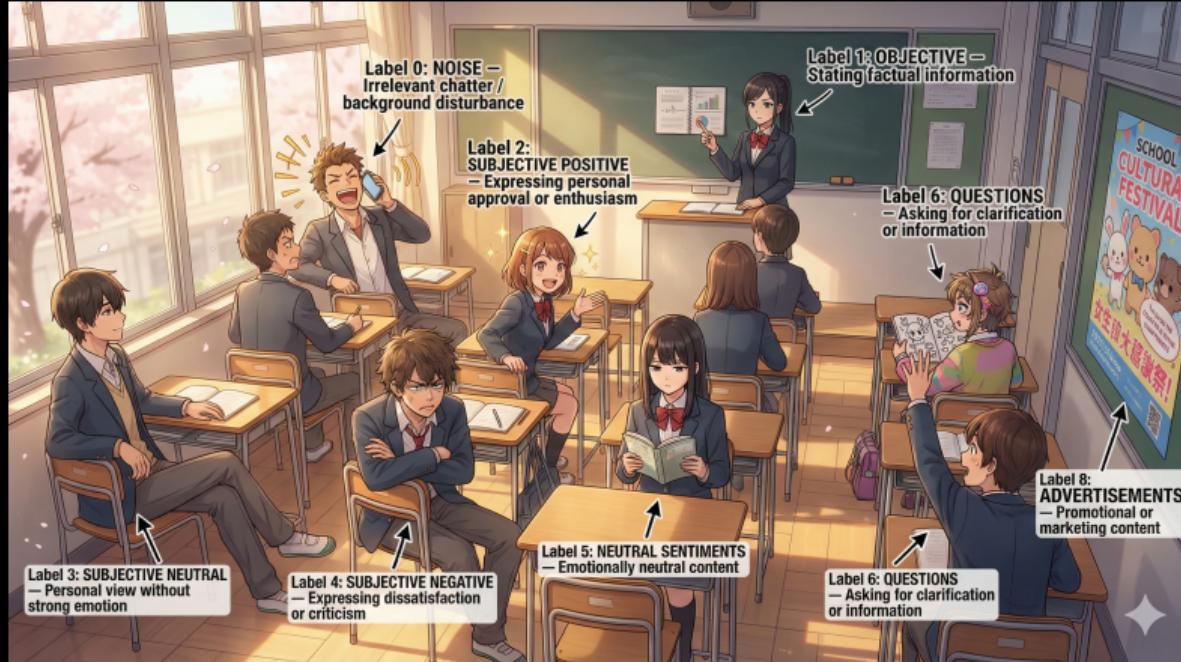
1. **Objective (Level 1):**
"Binance is reportedly building a \$1 billion insurance fund to counter crypto hacks."
2. **Subjective → Negative (complaint):**
"Do yourself a favor and don't use Kucoin... They are not reliable... my coins are now lost..."
3. **Subjective → Neutral → Question (intent):**
"How does TDS work? ... they give our 1% TAX with our PAN... can someone explain?"

Hierarchical Structure of Opinions



- ◊ **Level 1:** Coarse-grained (Noise, Objective, Subjective)
- ◊ **Level 2:** Sentiment refinement (Neutral, Negative, Positive)
- ◊ **Level 3:** Fine-grained neutral categories

Label Visualization



Approach 1: Classification Fine-tuning

- ◊ Custom classification head
- ◊ Parameter-efficient fine-tuning
- ◊ Weighted cross-entropy loss

Approach 2: Instruction Fine-tuning

- ◊ Prompt-response format
- ◊ Next-token prediction
- ◊ Masked loss on answer tokens

Base Model: [Gemma3-1B](#) — Compact 1-billion parameter LLM

3. Problem Statement

The 8-Class Classification Task

Hierarchical → Flat Mapping:

ID	Category
0	NOISE
1	OBJECTIVE
2	SUBJECTIVE_POSITIVE
3	SUBJECTIVE_NEGATIVE
4	NEUTRAL_SENTIMENTS
5	QUESTIONS
6	MISCELLANEOUS
7	ADVERTISEMENTS

Key Challenges

- ◊ Severe **class imbalance**
- ◊ **Semantic overlap** among categories
- ◊ **Noisy** social media text
- ◊ **Hierarchical dependencies**

4. Data Preparation

4-Phase Data Cleaning Workflow

Phase 1: Structural Integrity

- ◊ Schema alignment across sources
- ◊ Null entry elimination
- ◊ Token-based duplicate detection
- ◊ Corpus integrity verification

Phase 3: Text Normalization

- ◊ Case normalization (lowercase)
- ◊ Domain-specific token preservation
- ◊ Final text validation

Phase 2: Content Cleaning

- ◊ URL/hyperlink removal
- ◊ Mention/hashtag filtering
- ◊ Emoji normalization
- ◊ Whitespace compacting

Phase 4: Label Formatting

- ◊ Hierarchical label consolidation
- ◊ Numeric label encoding (0–8)
- ◊ Human-readable mapping

Dataset Statistics

Reddit Dataset

- ◇ 5,000 samples
- ◇ Mean: 186 tokens
- ◇ Max: 15,535 tokens

Top classes:

- ◇ Noise: 645
- ◇ Objective: 503
- ◇ Neutral: 476

Twitter Dataset

- ◇ 4,987 samples
- ◇ Mean: 46 tokens
- ◇ Max: 151 tokens

Top classes:

- ◇ Objective: 1,700
- ◇ Noise: 1,338
- ◇ Positive: 268

YouTube Dataset

- ◇ 5,000 samples
- ◇ Mean: 38 tokens
- ◇ Max: 1,128 tokens

Top classes:

- ◇ Negative: 1,574
- ◇ Neutral: 1,391
- ◇ Question: 1,000

Class Imbalance

Significant variation across platforms — *Advertisements* has only 1 sample in YouTube!

5. Model Architecture

Architecture Pipeline:

1. Token Embedding

$H_0 \leftarrow \text{EmbeddingLookup}(X_{\text{tokens}})$

2. Transformer Backbone

L transformer layers with Gemma3RMSNorm

3. Final Normalization

$H_{\text{final}} \leftarrow \text{RMSNorm}(H_L)$

4. Last Token Extraction

$h_{\text{last}} \leftarrow H_{\text{final}}[:, -1, :]$

5. Classification Head

LayerNorm → Dropout → Linear(8)

Custom Classification Head

```
Sequential(  
    LayerNorm(1152)  
    Dropout(p=0.4)  
    Linear(1152 → 8)  
)
```

Key Design

Replace original LM head
(1152 → 262144)
with classification head
(1152 → 8)

Trainable Components

- ◊ Last transformer block
- ◊ Final LayerNorm
- ◊ LM head
- ◊ Custom classification head

Trainable Components

- ◊ Last transformer block
- ◊ Final LayerNorm
- ◊ LM head
- ◊ Custom classification head

Frozen Components

- ◊ All earlier transformer blocks
- ◊ Preserves pretrained linguistic knowledge

Trainable Components

- ◊ Last transformer block
- ◊ Final LayerNorm
- ◊ LM head
- ◊ Custom classification head

Frozen Components

- ◊ All earlier transformer blocks
- ◊ Preserves pretrained linguistic knowledge

Benefits:

- ◊ Reduces trainable parameters significantly
- ◊ Prevents overfitting on small datasets
- ◊ Enables training on **single 24GB GPU**

6. Training Methodology

Efficiency Techniques

Memory Optimization

- ◊ **4-bit Quantization (NF4)**
- ◊ Reduces memory by $\sim 75\%$
- ◊ Uses bitsandbytes backend

Optimizer: AdamW

Learning rate	5×10^{-5}
Weight decay	0.1
Scheduler	Cosine + warmup
Warmup ratio	15%

Gradient Stability

- ◊ Gradient clipping (max norm = 1.0)
- ◊ bfloat16 precision
- ◊ Prevents NaN values

Training Config

- ◊ 3 epochs
- ◊ 70/20/10 train/val/test split
- ◊ Batch size: 4

Weighted Cross-Entropy Loss:

Class weights based on inverse frequency:

$$\text{weights} = [225, 175, 80, 130, 175, 900, 70, 30]$$

$$\text{normalized} = \frac{\text{weights}}{\sum \text{weights}} \times 8$$

Effect:

- ◊ Higher penalty for minority classes
- ◊ Forces model to learn rare categories

Without Weighting

- ◊ Model overfits to *Noise and Objective*
- ◊ Near-zero F1 for minorities
- ◊ Accuracy appears high but misleading

7. Instruction Fine-Tuning

Instruction Fine-Tuning Approach

Key Steps:

1. Tokenization of Prompt-Response Pairs

Concatenate prompt + input text

2. Sequence Shifting

$Y = [x_1, x_2, \dots, x_N]$ from $X = [x_0, x_1, \dots, x_{N-1}]$

3. Masking Non-Label Tokens

Set prompt tokens to -100 (ignored)

4. Label Alignment Verification

Ensure exact match of label tokens

Model Used: Gemma3:27B

Objective

- ◊ Next-token prediction
- ◊ Loss only on answer tokens
- ◊ Preserves autoregressive nature

Benefit

Better alignment with instruction-based reasoning

8. Results

Experimental Comparison: Classification vs. Instruction Tuning

Run 1: Classification Fine-tuning

Method: Standard supervised learning (Flat 8-class head).

Reddit (500 samples):

- ◊ **Subjective:** 41.4% (Dominated by opinion)
- ◊ **Objective:** 31.6%
- ◊ **Noise:** 27.0%

Twitter (500 samples):

- ◊ **Noise:** 64.4% (Model collapsed to majority)
- ◊ **Objective:** 18.2%
- ◊ **Subjective:** 17.4%

Run 2: Instruction Fine-tuning

Method: Prompt-response (Next-token prediction).

Reddit (500 samples):

- ◊ **Objective:** 44.6% (Better factual detection)
- ◊ **Noise:** 28.2%
- ◊ **Subjective:** 27.2%

Twitter (500 samples):

- ◊ **Noise:** 46.6% (Significant reduction)
- ◊ **Subjective:** 35.6% (Better opinion recovery)
- ◊ **Objective:** 17.8%

Impact on Minority Classes (QnA Dataset)

The Challenge

In the QnA dataset, **Class 1 (Relevant/Question)** is a severe minority compared to **Class 0 (Noise/Irrelevant)**.

Category	Run 1 (Class. Head)	Run 2 (Instruct.)
Class 0 (Majority)	99.75%	91.11%
Class 1 (Minority)	0.25%	8.89%

Improvement Factor

35 ×

Improvement in minority class recognition.

Why Instruction Tuning Won?

- ◇ **Reduced Noise Dominance:** Run 2 successfully shifted predictions away from the "safe" majority class (Noise).
- ◇ **Minority Coverage:** Significant gains in *Questions* and *Advertisements* categories.
- ◇ **Alignment:** The prompt-response format aligns better with the model's pre-trained reasoning capabilities.

Remaining Challenges

- ◇ **Noise Persistence:** Even in Run 2, noise remains a high percentage for Twitter (46.6%), due to the messy nature of the platform's data.
- ◇ **Subjective/Objective Flip:** In Reddit, Run 2 favored Objective labels more than Run 1, altering the distribution profile.

Conclusion: Instruction fine-tuning offers superior robustness for imbalanced, hierarchical datasets compared to standard classification heads.

9. Challenges & Solutions

Gradient Explosion

The Problem:

- ◊ Training destabilized early on.
- ◊ Large gradients caused weight updates to spiral.
- ◊ Resulted in **NaN values** in output tensors.

The Solution:

1. **Gradient Clipping:** Capped max norm at 1.0 to constrain updates.
2. **Precision Shift:** Switched from `float32` to `bfloat16`.
3. *Result:* Prevents numerical overflow while maintaining dynamic range.

Memory Overflow (OOM)

The Problem:

- ◊ 24GB VRAM limit on single GPU.
- ◊ Quadratic complexity of attention mechanism.
- ◊ `float32` weights exceeded capacity immediately.

The Solution:

1. **4-bit Quantization:** Used NF4 (`NormalFloat4`) via `bitsandbytes` ($\sim 75\%$ reduction).
2. **Batch Size Reduction:** Reduced training batch size from $8 \rightarrow 4$.
3. *Result:* Enabled fine-tuning of 1B parameters on consumer hardware.

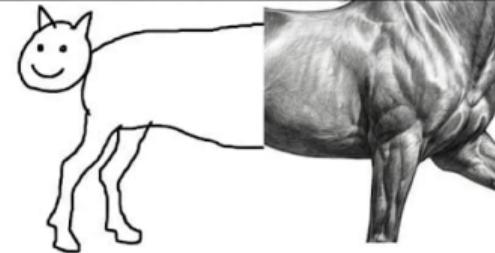
Gemma-1B Backbone (Frozen Layers)

A pre-trained linguistic powerhouse. Understands English perfectly. Ready for action.



Our New Classification Head & Fine-Tuning Process

Attaching a simple MLP. But wait...
Gradient Explosions! Memory Overflow!
Catastrophic Forgetting! The struggle
is real.



The Challenge: Performing '**'Brain Surgery'** without Breaking the Patient. We had to use **PEFT**, **bfloat16**, and gradient **clipping** just to get the new head to talk to the old body!

Overfitting to Majority Classes

The Problem:

- ◊ Model achieved high surface-level accuracy.
- ◊ *Reality*: It was predicting "Noise" or "Objective" for everything.
- ◊ **F1-Score** ≈ 0 for minority classes (Questions, Ads).

The Solution:

- ◊ **Weighted Cross-Entropy Loss**:
- ◊ Calculated inverse frequency weights:

$$w_c = \frac{N}{N_c}$$

- ◊ Penalized misclassification of rare classes heavily.

Stagnant Convergence

The Problem:

- ◊ Accuracy plateaued at ~50% during mid-training.
- ◊ Model unable to escape local minima.

The Solution:

- ◊ **Architecture Tuning**: Added extra LayerNorm before the head.
- ◊ **Regularization**: Introduced Dropout ($p=0.4$) to prevent neuron co-adaptation.
- ◊ **Activation**: Switched to GELU for smoother gradients.

10. Conclusion

- ◇ Adapted **Gemma-1B** for **8-class hierarchical opinion classification**
- ◇ Implemented **parameter-efficient fine-tuning**:
 - Custom classification head
 - Selective layer unfreezing
 - 4-bit quantization for memory efficiency
- ◇ Compared two approaches:
 - **Classification fine-tuning**: Direct supervised learning
 - **Instruction fine-tuning**: Better minority class coverage
- ◇ Addressed **class imbalance** via weighted loss functions
- ◇ Demonstrated effectiveness on **Reddit, Twitter, YouTube, QnA** datasets

Extensions

- ◊ Cross-domain generalization
- ◊ Cross-lingual adaptation
- ◊ Hierarchical label dependency modeling

Improvements

- ◊ Robustness to noisy data
- ◊ Temporal pattern analysis
- ◊ Refined evaluation metrics

Code Available:

https://github.com/Shuvam-Banerji-Seal/FIRE_2025_Hierarchical_Opinion_Classification_using_Large_Language_Models

11. Acknowledgments

We thank **Dr. Dwaipayan Roy**

Department of Computational and Data Sciences (CDS)
Indian Institute of Science Education and Research, Kolkata

for providing the computing resources for our fine-tuning runs.

Thank You!

Questions?